

handoff

How Handoff uses AWS

Build It: the open-source **Strands Agents SDK**. Ship It: **ten AWS services**, live in `ap-northeast-2`. Every claim in this deck sits next to a capture from the real account or the real terminal.

First Commit · Bharat Builds Tour

Ship It track

Verified live 20 Sept 2026

What the project does

Describe a recurring chore in one sentence — spoken or typed — and Handoff runs it on a schedule on AWS. It handles what it can judge confidently, and stops to ask you, once, only about the things it genuinely cannot call. Your answer becomes a rule, so it asks less every week.

Workflow	Wakes on	Tools	Asks you when
Morning Inbox Triage	weekdays 08:00	Gmail · Linear · Slack	a message is genuinely unclear
PR Review Triage	weekdays 09:00	GitHub · Slack	a PR touches auth, billing, migrations
Competitor Pricing Watch	Mondays 09:00	web · browser · Slack	the post reads as a strategic call
Slack Channel Digest	weekdays 18:00	Slack · Linear	you are mentioned but the ask is unclear
Meeting Follow-up	a meeting ending	Gmail · Linear	mail would leave the company

Five templates in `src/handoff/workflows/` ; anything else is a sentence to the Builder agent. One gate, one batch question, one memory — for all of them.

The short answer

BUILD IT • OPEN SOURCE

Strands Agents SDK 1.55

Every agent, the run graph, the human-in-the-loop gate, sessions, MCP tool clients and tracing are Strands primitives — 16 custom `@tool` s on top.


Also from AWS's open-source shelf: `bedrock-agentcore` (runtime + browser SDK), `amazon-transcribe` (streaming SDK), `boto3` .

The whole product runs **offline with no AWS account**: scripted model, JSON stores, `make demo` .

SHIP IT • AWS SERVICES

Service	Job
Amazon Bedrock (Nova Pro / Lite)	reasoning for every agent
Bedrock AgentCore Runtime	runs the graph, serverless
Bedrock AgentCore Memory	learned rules across runs
EventBridge Scheduler	cron triggers
AWS Lambda	tick → runtime bridge
Amazon DynamoDB	all state, one table
Amazon Transcribe	hears you, streaming
Amazon Polly	speaks back
Amazon ECR	the runtime's arm64 image
IAM + CloudWatch Logs	three scoped roles · logs

Architecture



Official AWS Architecture Icons · source `docs/architecture.drawio` , generated by `scripts/make_drawio_architecture.py`

Proof zero: it is live today

```
→ handoff (main) $ python scripts/aws_live_check.py
✓ IAM identity          arn:aws:iam::9410*****:user/handoff-deploy
✓ AgentCore Runtime     handoff READY
✓ AgentCore Memory      handoff_preferences-aWqi0xD7h0 ACTIVE
✓ DynamoDB              handoff ACTIVE PAY_PER_REQUEST keys=pk+sk items~680
✓ Lambda                handoff-tick python3.12 Active
✓ EventBridge Scheduler handoff-inbox-triage-morning ENABLED, handoff-competitor-pricing-watch ENABLED,
handoff-slack-channel-digest ENABLED, handoff-pr-review-triage ENABLED
✓ ECR                   handoff (3 images)
✓ Bedrock (Nova profiles) apac.amazon.nova-lite-v1:0, apac.amazon.nova-micro-v1:0, apac.amazon.nova-pro-v1:0,
global.amazon.nova-2-lite-v1:0
✓ Polly neural voices    13 en-US voices
✓ IAM roles              handoff-runtime, handoff-scheduler, handoff-tick

10/10 live in ap-northeast-2 · 2026-09-20 19:16
```

Captured 20 September 2026, 19:16 IST. `scripts/aws_live_check.py` makes read-only describe/list calls against the real account (id masked). Ten of ten resources live. The DynamoDB table held 67 items when first captured on 13 Sept and ~680 now — the schedules have been firing all week.


PART ONE

Build It

The open-source AWS stack: **Strands Agents SDK**, plus the `bedrock-agentcore` and `amazon-transcribe` SDKs. No account, no card, no bill — the whole loop runs on a laptop.

Strands Agents SDK — every feature doing a real job

Strands feature	Where it earns its place
Agent	builder, voice assistant, executor, completer, learner, chat
@tool	16 custom tools — <code>fetch_unread_emails</code> ... <code>check_competitor_pricing</code>
Graph + GraphBuilder + conditional edges	trigger → executor → completer; an empty trigger never costs a model call
BeforeToolCallEvent + event.interrupt()	the human-in-the-loop gate
Graph.serialize_state	a paused run survives a restart
Invocation / tool hooks	the narrator that draws the live run graph
SessionRepository	terminal and browser share one chat
MCPClient (stdio + HTTP)	Gmail, Linear, Slack, GitHub, Notion, Airtable, web
AgentTool / OpenAIModel subclasses	host embedding · tool-JSON repair
OpenTelemetry tracing	<code>config.configure_observability()</code>



```
→ handoff (main) $ python -c 'import strands; ...'
strands-agents 1.55.1
strands-agents-tools 0.8.8
bedrock-agentcore 1.23.0

14 from strands import tool
5 from strands import Agent
2 from strands.tools.mcp import MCPClient
1 from strands.types.tools import AgentTool, ToolResult, ToolSpec, ToolUse
1 from strands.types.session import Session, SessionAgent, SessionMessage, SessionType
1 from strands.types.events import ToolResultEvent
1 from strands.telemetry import StrandsTelemetry
1 from strands.session.session_repository import SessionRepository
1 from strands.session.repository_session_manager import RepositorySessionManager
1 from strands.multiagent import GraphBuilder
1 from strands.models.openai import OpenAIModel
1 from strands.models.model import Model
1 from strands.models.bedrock import BedrockModel
1 from strands.models.anthropic import AnthropicModel
```

Installed versions and every Strands import in the codebase
— `strands-agents 1.55.1` .

The gate — one Strands line the product is built around

```
# src/handoff/graph/hooks/hitl.py
# First pass: raises InterruptException and
# suspends the agent loop — before the tool runs.
# After resume: returns the human's answer.
response = event.interrupt(
    interrupt_name,
    reason=payload.model_dump(mode="json"),
)
```

- A `BeforeToolCallEvent` hook sits in front of the **only tool that changes anything** in the outside world.
- Below the confidence threshold it **defers**; `finish_batch` then raises **one** interrupt for the whole pass — one screen, not one interruption per item.
- The answer becomes a narrow rule (sender, domain, or two keywords).



```
→ handoff (main) $ handoff decide int_3990c60b15be archive --note 'cold outreach'
[handoff:slack] Please review the emails from Vendor Corp and HR for the strategic partnership opportunity and upda
[handoff:slack] Morning triage complete. 6 handled on my own, 2 decided by you, 0 handled from rules you set earlie
[handoff:slack] Please review the emails from Vendor Corp and HR for the strategic partnership opportunity and upda
✓ archive recorded — run completed
learned: a new rule
```


2 learned rules					
id	action	rule	matches	confidence	applied
pref_e9b29952645a	archive	Archive emails related to cold outreach as they are not worth a ticket	cfo@company.com	80%	0
pref_a8a635211554	archive	Archive any cold outreach from partnerships@vendor.com	partnerships@vendor.com	90%	0

`handoff decide` on a real run: the human's answer resumes the same Strands graph, and the learner writes the rule.

Build It means it runs with nothing

```
git clone https://github.com/ZINKUNO/Handoff-AWS && cd Handoff-AWS
python3.12 -m venv .venv && source .venv/bin/activate
pip install -e ".[dev,bedrock,desktop,web,voice]"
make demo          # scripted model · JSON stores · no keys, no account
```

- Every cloud store has a **local JSON twin**: DynamoDB ↔ `.handoff-state/*.json`, AgentCore Memory ↔ local preferences.
- `HANDOFF_FAKE_MODEL=true` swaps Bedrock for a deterministic scripted model — that is what the test suite runs on.
- `--state-dir` always means local storage, so a scratch run **can never touch the live table**.

A terminal window with a dark background and light green text. It shows the command `pytest tests/ -q -p no:warnings && ruff check src tests scripts infra` being executed. The output shows progress bars for different test categories: 27%, 54%, 81%, and 100%. It then reports `266 passed in 105.48s (0:01:45)` and `All checks passed!`.

```
→ handoff (main) $ pytest tests/ -q -p no:warnings && ruff check src tests scripts infra
..... [ 27%]
..... [ 54%]
..... [ 81%]
..... [100%]
266 passed in 105.48s (0:01:45)
All checks passed!
```

266 tests, lint clean — re-run on 20 Sept 2026.

`tests/test_hitl_gate.py` drives the gate inside a real Strands agent: an unsure call stops the loop **and the tool never runs**.

PART TWO

Ship It

Ten AWS services, deployed by boto3 scripts in `infra/` — no console clicking. One slide per service: what it does, the code that calls it, and the capture that proves it.

Amazon Bedrock — Nova Pro and Nova Lite

Job: the reasoning behind all six agents, through cross-region inference profiles.

```
# src/handoff/config.py
def bedrock_profile(model_id, region=None):
    """Resolve a bare Bedrock model id to the
    inference profile for a region."""
    if model_id.startswith(_GEO_PREFIXES):
        return model_id
    geo = _GEO_BY_REGION.get(
        (region or AWS_REGION).split("-")[0], "")
    return f"{geo}-{model_id}"
...
return BedrockModel(model_id=model_id,
                     region_name=AWS_REGION)
```

Decision: Nova Pro for the demo take, **Nova Lite for tests** — a spoken turn on Lite costs under a tenth of a cent.

```
→ handoff (main) $ aws bedrock list-inference-profiles && aws polly describe-voices && handoff doctor speech bedr
```

ListInferenceProfiles	
profile	status
apac.amazon.nova-micro-v1:0	ACTIVE
apac.amazon.nova-lite-v1:0	ACTIVE
apac.amazon.nova-pro-v1:0	ACTIVE
global.amazon.nova-2-lite-v1:0	ACTIVE

Polly neural voices, ap-northeast-2:
Danielle Gregory Joanna Ruth Kevin Salli Matthew Kimberly Kendra Justin Joey Ivy Stephen

Handoff doctor — provider: bedrock

PASS Speech Polly Matthew (neural) + Transcribe streaming in ap-northeast-2
PASS AWS Bedrock apac.amazon.nova-pro-v1:0 in ap-northeast-2

2 working, 0 broken, 0 not configured

`aws bedrock list-inference-profiles` — four Nova profiles ACTIVE in ap-northeast-2 ; `handoff doctor` then makes a real Bedrock call.


Bedrock AgentCore Runtime — where the graph runs

Job: serverless background execution. The Strands graph ships as an **arm64 container**; each run gets its own session; long-running invocations are fine.

```
# src/handoff/app.py
from bedrock_agentcore import BedrockAgentCoreApp
app = BedrockAgentCoreApp()

@app.entrypoint
def handler(event, context=None):
    return handle(event, context)
```

```
python infra/deploy_agentcore.py --check # what is missing
python infra/deploy_agentcore.py         # build → ECR → Runtime
```



```
→ handoff (main) $ aws bedrock-agentcore-control list-agent-runtimes / list-memories
```

ListAgentRuntimes			
name	status	updated	version
handoff	READY	2026-09-13T08:11:37.761866+00:00	1

ListMemories	
created	id
2026-09-13T13:08:02.267000+05:30	handoff_preferences-awqi0xD7h0
	status: ACTIVE

`aws bedrock-agentcore-control list-agent-runtimes` → **handoff** · **READY**. Still READY in today's live check.

Bedrock AgentCore Memory — it learns the person

Job: every human decision becomes a narrow rule, stored in AgentCore Memory and recalled at the start of the next run — so it asks less every week.

```
# src/handoff/memory/store.py
client.create_event(
    memoryId=config.AGENTCORE_MEMORY_ID,
    actorId=pref.preference_key,
    sessionId=pref.source_interrupt_id,
    payload=[{"conversational": {...}}],
)
...
client.retrieve_memory_records(
    memoryId=config.AGENTCORE_MEMORY_ID,
    namespace=f"/preferences/{preference_key}",
    searchCriteria={"searchQuery": query, "topK": 5},
)
```

```
→ handoff (main) $ handoff doctor | grep -E "provider|Bedrock|Speech|DynamoDB|AgentCore"
Handoff doctor - provider: bedrock
PASS AWS Bedrock      apac.amazon.nova-pro-v1:0 in ap-northeast-2
PASS Speech           Polly Matthew (neural) + Transcribe streaming in ap-northeast-2
PASS DynamoDB         handoff (680 items)
PASS AgentCore Memory handoff_preferences-aWqi0xD7h0 (ACTIVE)

# the AWS checks only - each line is a real call made on 2026-09-20
```


handoff doctor , 20 Sept 2026 — the AWS checks.
Memory store handoff_preferences-aWqi0xD7h0 ACTIVE;
every line is a real call, not a config lint.

EventBridge Scheduler → Lambda — the cron that wakes it

Job: one schedule per active workflow. Scheduler **cannot target AgentCore directly**, so a twelve-line Lambda forwards the tick, unchanged.

```
# infra/lambda_setup.py – the whole function
client = boto3.client("bedrock-agentcore")

def handler(event, context):
    response = client.invoke_agent_runtime(
        agentRuntimeArn=RUNTIME_ARN,
        runtimeSessionId=uuid.uuid4().hex + ...,
        payload=json.dumps(event).encode(),
        contentType="application/json",
    )
    body = response["response"].read()
    return {"ok": True, "bytes": len(body)}
```



```
→ handoff (main) $ aws scheduler list-schedules && aws lambda list-functions && aws ecr describe-repositories
```

ListSchedules		
name	state	target
handoff-inbox-triage-morning	ENABLED	arn:aws:lambda:ap-northeast-2:9410*****:function:handoff-tick
handoff-competitor-pricing-watch	ENABLED	arn:aws:lambda:ap-northeast-2:9410*****:function:handoff-tick
handoff-slack-channel-digest	ENABLED	arn:aws:lambda:ap-northeast-2:9410*****:function:handoff-tick
handoff-pr-review-triage	ENABLED	arn:aws:lambda:ap-northeast-2:9410*****:function:handoff-tick

ListFunctions			
arch	name	runtime	updated
x86_64	handoff-tick	python3.12	2026-09-13T08:20:37.341+0000

DescribeRepositories	
repo	uri
handoff	9410*****.dkr.ecr.ap-northeast-2.amazonaws.com/handoff


Four schedules ENABLED, all targeting `handoff-tick` ; the function; the ECR repo. One real `aws` CLI session, account id masked.

Amazon DynamoDB — all state, one table

Job: workflow configs, runs, interrupts, chat sessions, usage and the audit trail — **one table**, `pk` = collection, `sk` = id, **pay-per-request**.

```
# src/handoff/store.py — DynamoCollection
kwargs = {"KeyConditionExpression":
          Key("pk").eq(self.collection)}
response = self._table.query(**kwargs)
...
self._table.put_item(
    Item={**payload, **self._key(payload[key_field])})
```

Decisions: single-table so a page costs **one Query per collection**; a 3-second read cache because reads from India to Seoul are ~150 ms each; on-demand billing because an idle agent should cost nothing.



```
→ handoff (main) $ aws dynamodb describe-table --table-name handoff
{
  "name": "handoff",
  "status": "ACTIVE",
  "items": 67,
  "billing": "PAY_PER_REQUEST",
  "keys": [
    {
      "AttributeName": "pk",
      "KeyType": "HASH"
    },
    {
      "AttributeName": "sk",
      "KeyType": "RANGE"
    }
  ]
}
```


`aws dynamodb describe-table`, 13 Sept: ACTIVE, PAY_PER_REQUEST, 67 items. Today: ~680.

Amazon Transcribe + Amazon Polly — voice in, voice out

Job: you say the chore. Audio streams to **Transcribe** over a WebSocket while you are still talking (partials fill the caption live); **Polly** answers.

```
# src/handoff/speech/aws.py
stream = await client.start_stream_transcription(
    language_code=language,
    media_sample_rate_hz=rate,
    media_encoding="pcm",
    enable_partial_results_stabilization=True,
    partial_results_stability="medium",
)
...
client.synthesize_speech(Text=text[:2900],
    VoiceId=voice, Engine=engine, OutputFormat="mp3")
```

Round-trip check: Polly said a sentence, Transcribe returned it word for word. The demo film's narration is Polly's generative voice too.




Talk — the orb. The caption fills in from Transcribe partials; four tool calls later the workflow is saved and running.

ECR, IAM and CloudWatch — the plumbing

- **Amazon ECR** — repo `handoff` holds the runtime's arm64 image (3 images today).
- **IAM** — three roles, each with one job:
 - `handoff-runtime` (the container), `handoff-tick` (Lambda → `InvokeAgentRuntime`), `handoff-scheduler` (Scheduler → that one function). A dedicated `handoff-deploy` user does the deploying.
- **CloudWatch Logs** — `/aws/bedrock-agentcore/runtimes/handoff-...` and `/aws/lambda/handoff-tick` both hold data from scheduled runs.

```
python infra/iam_setup.py          # scheduler role
python infra/lambda_setup.py --runtime-arn ...
python infra/eventbridge_setup.py \
  --target-arn <lambda> --role-arn <role>
```



```
→ handoff (main) $ aws sts get-caller-identity && aws configure get region
{
  "UserId": "AIDA5WHKRR2E02AD2GB4N",
  "Account": "9410*****",
  "Arn": "arn:aws:iam::9410*****:user/handoff-deploy"
}

region: ap-northeast-2
```

`aws sts get-caller-identity` — the deploy identity and region. Account id masked in every capture.

A real run on Bedrock — inbox triage

`handoff run --watch` streams the Strands hook events as they happen:

- **trigger** checks there is work
 - **memory** applies a learned rule
 - **Gmail** reads eight messages
 - the **executor** handles seven alone
 - **one** is set aside for a human
- 8 items in · 7 handled alone · 1 question.

Captured on Bedrock Nova, mid-September 2026.

```
→ handoff (main) $ handoff run inbox-triage-morning --watch
18:20:46 fetched    Fetched 8 unread messages (synthetic inbox)
18:20:46 tool_end    fetch_unread_emails ok
18:20:51 tool_start  calling submit_action
18:20:51 acted      Filed a ticket for "Need API review by Thursday" - 90% sure
18:20:51 tool_start  calling submit_action
18:20:51 acted      Archived "TC Daily: AI Agents Are Eating Software" - 80% sure
18:20:51 tool_start  calling submit_action
18:20:51 acted      Drafted a reply to "Quick sync on Q3 targets" - 90% sure
18:20:51 tool_start  calling submit_action
18:20:51 deferred    Not sure about "Strategic partnership opportunity - time sensitive" (60%) - setting it aside fo
18:20:51 tool_start  calling submit_action
18:20:51 acted      Left alone "Updated PTO policy - action may be required" - 70% sure
18:20:51 tool_start  calling submit_action
18:20:51 acted      Filed a ticket for "Re: Sprint retro action items" - 90% sure
18:20:51 tool_start  calling submit_action
18:20:51 acted      Filed a ticket for "[handoff/handoff] Dependabot alert: lodash prototype pollution" - 90% sure
18:20:51 tool_start  calling submit_action
18:20:51 tool_start  calling finish_batch
18:20:51 memory      Archived "FYI: Budget freeze through end of month" (a rule you set) - 80% sure
18:20:51 tool_end    submit_action ok
18:20:51 tool_end    submit_action ok
18:20:51 tool_end    submit_action ok
18:20:51 tool_end    submit_action ok
18:20:51 tool_end    submit_action ok
18:20:51 tool_end    submit_action ok
18:20:51 tool_end    submit_action ok
18:20:51 node_end    executor done
[handoff:console] Handoff needs your input
Item: Strategic partnership opportunity - time sensitive - from Unknown - Vendor Corp
From: partnerships@vendor.com
Why: Confidence (60%) is below the 70% threshold
My best guess: skip (60% confidence)
Decide: http://localhost:8000/decide/int_3990c60b15be
18:20:51 asked      Set aside 1 item for you
run      run_826bda5c0873
status   waiting_on_human
handled alone 6
decided by you 0
from your rules 1
summary     Paused on 1 decision that need you.
! 1 decision(s) need you - 'handoff pending'
```

A second real run — nothing to do with email

```
→ handoff (main) $ handoff run competitor-pricing-watch --watch
16:31:35 started Starting Weekly Competitor Pricing Watch (manual) on apac.amazon.nova-pro-v1:0
16:31:36 node_start trigger starting
16:31:43 tool_start calling check_trigger
16:31:43 tool_end check_trigger ok
16:31:44 tool_start calling check_trigger
16:31:44 tool_end check_trigger ok
16:31:46 node_end trigger done
16:31:46 node_start executor starting
16:31:47 tool_start calling recall_preferences
16:31:47 tool_end recall_preferences ok
16:31:48 tool_start calling check_competitor_pricing
16:31:48 tool_end check_competitor_pricing ok
16:31:50 node_end executor done
16:31:50 node_start completer starting
16:31:52 tool_start calling finalize_run
16:31:52 tool_start calling notify_user
16:31:53 note The pricing for the Enterprise tier increased from $99 to $129 per month. This change was handled
autonomously, but it reads as a strategic call, so I need your confirmation before posting it to #competitive-intel.
16:31:53 tool_end finalize_run ok
16:31:57 tool_end notify_user ok
16:31:59 node_end completer done
16:31:59 completed Done - 0 handled alone, 0 decided by you, 0 from your rules
run run_5fe989bb78d3
status completed
```

`competitor-pricing-watch` on Bedrock Nova Pro, 20 Sept 2026, against the sample pricing page in the repo. It calls its own tool, finds Enterprise moved \$99 → \$129, and **does not post** to #competitive-intel on its own — a strategic call, so it asks first.

Architecture and cost decisions

Decision	Why
AgentCore Runtime , not an always-on server	an agent that runs a few minutes a day should cost nothing the rest of it
EventBridge → Lambda → Runtime	Scheduler cannot target AgentCore; twelve lines bridge it, one permission each
One DynamoDB table, pay-per-request	no capacity to size, one Query per page, idle = free
Nova Pro to demo, Nova Lite to test	a spoken turn on Lite is under a tenth of a cent; a full demo take on Pro is cents
Conditional edge after the trigger	an empty inbox or an unchanged page never costs a model call
Local JSON twin for every store	development and the 266-test suite cost \$0 and need no account
Voice matched locally	"archive it", "leave it" resolve with no model round-trip

Transcribe streaming is \$0.024/min and Polly \$16 per million characters.

What fought back — what we learned

- **Bedrock inference profiles.** Model ids need a geography prefix (`us.` , `apac.` , `eu.`) that must match the *calling* region. Handoff now resolves a bare id against `AWS_REGION` .
- **Marketplace billing.** Anthropic models on Bedrock are sold through AWS Marketplace; an account that cannot complete that agreement gets `INVALID_PAYMENT_INSTRUMENT` in every region. Amazon's own Nova models just work — so Nova Pro became the default.
- **Scheduler ≠ AgentCore.** EventBridge Scheduler has no AgentCore target, and the old `agentcore` CLI "launches" nothing. The fix: `create_agent_runtime` through boto3, and a twelve-line Lambda bridge.
- **The Strands interrupt API is version-sensitive.** `event.interrupt(name, reason=...)` *returns* the human's answer on resume. Pinned to `strands-agents 1.55` , with tests that fail loudly if that contract moves.
- **Latency is a design input.** ~150 ms per DynamoDB read from India to Seoul turned into a single-table design and a short read cache.

What is not on AWS — stated plainly

- **The marketing site** (<https://handoff-aws.pages.dev>) is static files on Cloudflare Pages. The product — agents, schedule, state, memory, voice — is what runs on AWS.
- **Tool servers** (Gmail, Linear, Slack, GitHub, Notion, Airtable) are third-party MCP servers the agents call; they are integrations, not infrastructure.
- **What has real runs:** inbox triage against a live inbox, and the pricing watch on Bedrock against a sample pricing page. PR review triage, the Slack digest and meeting follow-up ship as templates on the same gate and need your own credentials.
- **Groq and the Anthropic API** exist as alternative model providers behind one switch (`HANDOFF_MODEL_PROVIDER`); the deployed runtime uses Bedrock.

Paste-ready answer for the form

What it does. Handoff turns one sentence — spoken or typed — into an autonomous workflow (inbox triage, PR review triage, competitor pricing watch, Slack digest, meeting follow-up). It runs on a schedule on AWS, acts alone where it is confident, asks one batched question where it is not, and turns each answer into a rule so it asks less every week.

Build It — AWS open source. Strands Agents SDK 1.55: `Agent` , `Graph` / `GraphBuilder` , `BeforeToolCallEvent` + `event.interrupt()` for the human-in-the-loop gate, `SessionRepository` , `MCPClient` , 16 `@tools` , `OpenTelemetry` tracing. Plus the open-source **bedrock-agentcore** and **amazon-transcribe** SDKs and **boto3**. Runs fully offline with no AWS account (`make demo` , 266 tests).

Ship It — AWS services. Amazon Bedrock (Nova Pro / Nova Lite) for reasoning · **Bedrock AgentCore Runtime** runs the graph as an arm64 container · **AgentCore Memory** stores learned rules · **EventBridge Scheduler** → **AWS Lambda** triggers runs · **Amazon DynamoDB** single table for all state · **Amazon Transcribe** (streaming) and **Amazon Polly** for voice · **Amazon ECR**, **IAM**, **CloudWatch Logs**. All live in `ap-northeast-2` , deployed by `boto3` scripts in `infra/` .

handoff

Describe it. Hand it off. It runs — on AWS.

github.com/ZINKUNO/Handoff-AWS

handoff-aws.pages.dev

Apache-2.0

Re-verify any slide yourself: `python scripts/aws_live_check.py · handoff doctor · make test`